



Leibniz
Universität
Hannover

Autoencoder

Julian Roth, Max Schröder,
Annika Heil, Yerso Checya Sinti

14.7.2020

Our Project



▶ Link



▶ Link



Table of Contents

1. Neural Networks

1.1 Activation Functions

1.2 Loss Functions

1.3 Optimizer

1.4 Code: Training a Classifier

2. Autoencoder

2.1 Principal Component Analysis

2.2 Autoencoder

2.3 Denoising Autoencoder

2.4 Variational Autoencoder

2.5 Generative Adversarial Network

3. Outlook



Table of Contents

1. Neural Networks

1.1 Activation Functions

1.2 Loss Functions

1.3 Optimizer

1.4 Code: Training a Classifier

2. Autoencoder

2.1 Principal Component Analysis

2.2 Autoencoder

2.3 Denoising Autoencoder

2.4 Variational Autoencoder

2.5 Generative Adversarial Network


3. Outlook

Activation Functions



- Sigmoid
- Tanh
- ReLU
- LeakyReLU
- Softmax

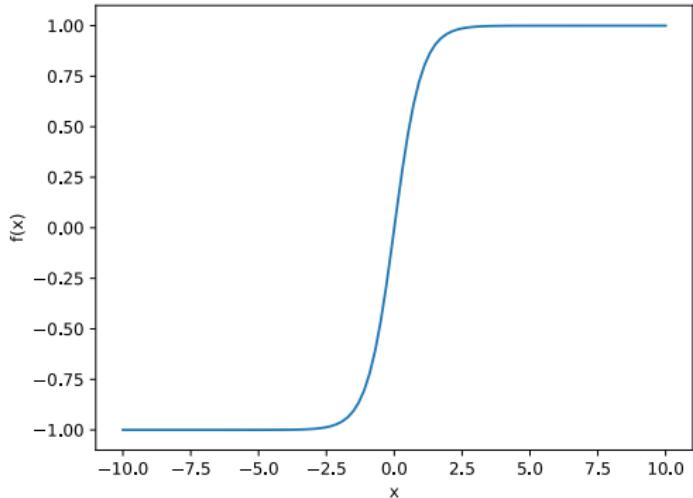
Sigmoid



- ⊕ biological neuron
- ⊕ smooth
- ⊖ vanishing gradient

$$\sigma(x) := \frac{1}{1 + e^{-x}}$$


Hyperbolic Tangent



- ⊕ similar to Sigmoid
- ⊕ used in LSTMs

$$\tanh(x) := \frac{\sinh(x)}{\cosh(x)}$$

Rectified Linear Unit



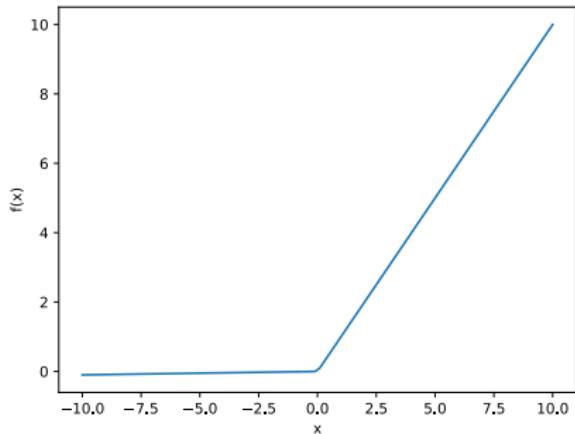
$$\text{ReLU}(x) := \max(0, x)$$

⊕ cheap

⊖ not differentiable in 0

$$\text{ELU}(x) := \begin{cases} x & \text{for } x \geq 0 \\ e^x - 1 & \text{for } x < 0 \end{cases}$$

Leaky Rectified Linear Unit



- + similar to ReLU
- + doesn't vanish for $x < 0$

$$\text{LeakyReLU}(x) := \max(\varepsilon x, x) \text{ with } \varepsilon \ll 1$$

Softmax



- probability distribution over classes
- output example:

$$\mathbb{P}(\text{"dog"}) = 80\%$$

$$\mathbb{P}(\text{"cat"}) = 20\%$$

$$\text{Softmax}(x_i) := \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

Loss Functions



- Mean Squared Error
- Crossentropy



Mean Squared Error (MSE)

- L^2 -error between the predictions of the neural network and the expected outputs
- used for regression tasks

$$\text{MSE}(X) := \frac{1}{2|X|} \sum_{x \in X} ||\text{NN}(x) - y(x)||^2$$

Crossentropy



- output of neural network should resemble probability distribution
- forces output of neural network to either be close to 0 or close to 1
- used for classification tasks

$$\text{Crossentropy}(X) := -\frac{1}{|X|} \sum_{x \in X} \left[y(x) \ln (\text{NN}(x)) + (1 - y(x)) \ln (1 - \text{NN}(x)) \right]$$



Adam Optimizer [8]

$$\begin{aligned}g_t &= \nabla_{\theta} \text{Loss}_t(\theta_{t-1}) \\m_t &= (1 - \beta_1)g_t + \beta_1 m_{t-1} \\v_t &= (1 - \beta_2)g_t^2 + \beta_2 v_{t-1} \\\hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \\\hat{v}_t &= \frac{v_t}{1 - \beta_2^t}\end{aligned}$$

$$\theta_{t+1} = \theta_t - \frac{\eta \cdot \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$



Justification of estimates [8]


Relation between $\mathbb{E}[v_t]$ and $\mathbb{E}[g_t^2]$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \Rightarrow v_t = (1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} g_i^2$$


$$\begin{aligned}\mathbb{E}[v_t] &= \mathbb{E} \left[(1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} g_i^2 \right] = \mathbb{E} \left[g_t^2 \right] (1 - \beta_2) \sum_{i=1}^t \beta_2^{t-i} + \zeta \\ &= \mathbb{E} \left[g_t^2 \right] (1 - \beta_2^t) + \zeta\end{aligned}$$

Similarly for $\mathbb{E}[m_t]$ and $\mathbb{E}[g_t]$.


Experimental results: MSE



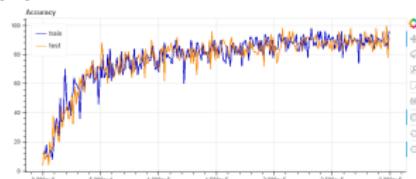
(a) Loss: SGD + MSE ($\eta = 1$)



(b) Accuracy: SGD + MSE (89%)




(c) Loss: Adam + MSE ($\eta = 0.1$)




(d) Accuracy: Adam + MSE (91%)


Experimental results: CE




(a) Loss: SGD + CE ($\eta = 1$)



(b) Accuracy: SGD + CE (95%)



(c) Loss: Adam + CE ($\eta = 0.1$)



(d) Accuracy: Adam + CE (98%)



Code: Training a Classifier

```
1 from nn import MLP
2 from layers import Dense
3 from activations import ReLU, Softmax
4 from loss import CrossEntropy
5 from dataset import Dataset
6 from optimizer import Adam
7
8 dataset = Dataset(name="mnist", train_size=60000, test_size=10000, batch_size=50)
9
10 classifier = MLP()
11 optimizer = Adam(learningRate = 0.1)
12
13 classifier.addLayer(
14     Dense(inputDim = 28 * 28, outputDim = 100, activation = ReLU(), optimizer = optimizer)
15 )
16 classifier.addLayer(
17     Dense(inputDim = 100, outputDim = 50, activation = ReLU(), optimizer = optimizer)
18 )
19 classifier.addLayer(
20     Dense(inputDim = 50, outputDim = 10, activation = Softmax(), optimizer = optimizer)
21 )
22
23 classifier.train(
24     dataset, loss = CrossEntropy(), epochs = 10,
25     metrics = ["train_loss", "test_loss", "train_accuracy", "test_accuracy"],
26     tensorboard = False, callbacks = {}
27 )
```



Table of Contents

1. Neural Networks

1.1 Activation Functions

1.2 Loss Functions

1.3 Optimizer

1.4 Code: Training a Classifier

2. Autoencoder

2.1 Principal Component Analysis

2.2 Autoencoder



2.3 Denoising Autoencoder

2.4 Variational Autoencoder

2.5 Generative Adversarial Network

3. Outlook

Principal Component Analysis (PCA)

 \mathbb{R}^2  \mathbb{R} 



PCA: Theory

Goal: Reduce $\{x^{(i)}\}_{i=1}^n \subset \mathbb{R}^d$ to k -dimensional data with $k \ll d$.

Want to learn more?



[Lecture 14](#)

[Machine Learning \(Stanford\) \[1\]](#)

Pre-processing: Zero out mean

```
μ = 1/n ∑i=0n x(i)
for i in range(1, n):
    x(i) -= μ
```



PCA: Theory [1, 12]

Let us revisit our model problem. We want to find a vector u^* s.t. most of the variation of the data is preserved, i.e.

$$u^* = \operatorname{argmax}_{u \in \mathbb{R}^2: \|u\|=1} \frac{1}{n} \sum_{i=1}^n (x^{(i)} \cdot u)^2. \quad (1)$$

Using the cosine similarity of two vectors and the definition of the cosine, we get that for $\|u\| = 1$, the projection of $x^{(i)}$ onto u has length $x^{(i)} \cdot u$. Thus we can see that we are trying to maximize the average squared length of our projection.



PCA: Theory [1, 12]

$$\begin{aligned}\frac{1}{n} \sum_{i=1}^n (x^{(i)} \cdot u)^2 &= \frac{1}{n} \sum_{i=1}^n (u^T x^{(i)}) (x^{(i) T} u) \\ &= u^T \underbrace{\left(\frac{1}{n} \sum_{i=1}^n x^{(i)} x^{(i) T} \right)}_{=: \Sigma} u\end{aligned}$$

Now we can rewrite (1) as

$$\operatorname{argmax} u^T \Sigma u \quad \text{s.t.} \quad u^T u = 1.$$



PCA: Theory [1, 12]

$$\operatorname{argmax} u^T \Sigma u \quad \text{s.t.} \quad u^T u = 1$$

Through the Lagrange formalism (\rightarrow lectures on optimization), we get

$$\mathcal{L}(u, \lambda) = u^T \Sigma u - \lambda (u^T u - 1).$$

Hence for a maximum holds

$$0 \stackrel{!}{=} \nabla_u \mathcal{L}(u, \lambda) = \Sigma u - \lambda u.$$

Therefore, we simply need to find the eigenvector u corresponding to the biggest eigenvalue of the covariance matrix Σ .



PCA: Theory [1, 12]

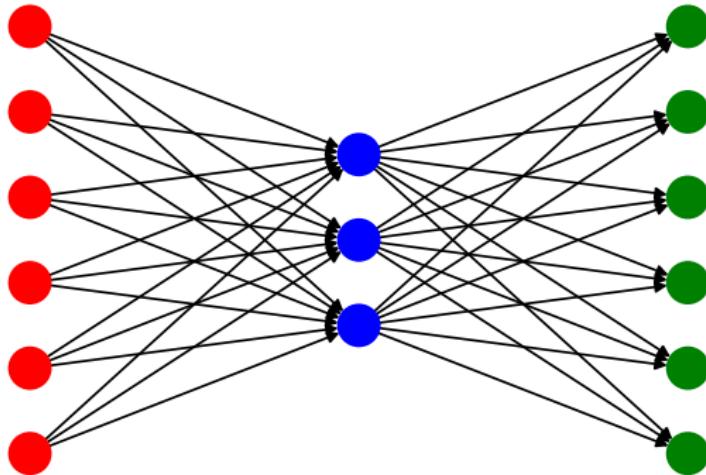
In general:

Find the eigenvectors u_1, \dots, u_k corresponding to the k biggest eigenvalues of the covariance matrix $\Sigma = XX^T$.



Alternative:

Find vectors u_1, \dots, u_k via singular value decomposition of X [12].

Autoencoder




Denoising Autoencoder






Application: TransCoder[9]

Cross-lingual Masked Language Model pretraining



Denoising auto-encoding



Back-translation




Figure from Marie-Anne Lachaux et al.



Variational Autoencoder Introduction

Goal: Generate new synthetic data

Idea: Construct autoencoder with low dimensional latent space


→ sample from low dimensional latent space, then use decoder to obtain new data

Observation: Generated data will mostly be random noise


New idea: Train autoencoder to organize data in latent space

Observation: Generated data is meaningful


AE



VAE



Latent Space Visualization for MNIST [2]




VAE Overview



Given: Dataset

Goal: Generate new meaningful data \hat{x} similar to original data

Data generation process we desire:




Note that \hat{x} is a sample and not calculated by a deterministic function.



Training Process

Train model to behave like an autoencoder: x input, \hat{x} output with $\hat{x} \approx x$
→ Ensures generation of meaningful data





Relation of Distributions

Assumptions for distributions:

$$F_Z \sim \mathcal{N}(0, I)$$

$$F_{X|Z} \sim \mathcal{N}(f(Z), cI), c > 0 \text{ fixed}, f \in \mathcal{F}$$

Bayes theorem/conditional densities:

$$f_{Z|X}(z|x) = \frac{f_{X,Z}(x,z)}{f_X(x)} = \frac{f_{X|Z}(x|z)f_Z(z)}{f_X(x)} = \frac{f_{X|Z}(x|z)f_Z(z)}{\int f_{X|Z}(x,y)f_Z(y)dy}$$

→ Can be used to determine $F_{Z|X}$ but inefficient to calculate
Solution: Use "variational inference" to approximate $f_{Z|X}$

Relative Entropy/Kullback-Leibler Divergence



Definition: For $P \ll Q$

$$H(P, Q) := E_P \left[\log \frac{dP}{dQ} \right]$$

Interpretation: "Distance" between probability distributions

Variational Inference



Let f and x be fixed (Remember: $F_{X|Z} \sim \mathcal{N}(f(Z), cI)$).

Approximate $f_{Z|x}(\cdot|x)$ by \tilde{f}_x with $\tilde{F}_x \sim \mathcal{N}(g(x), h(x))$ for $g \in \mathcal{G}, h \in \mathcal{H}$.

Goal: Find best approximation $\tilde{F}_x^* \sim \mathcal{N}(g^*(x), h^*(x))$

$$(g^*, h^*) = \arg \min_{(g,h) \in \mathcal{G} \times \mathcal{H}} H(\tilde{F}_x, F_{Z|x}(\cdot|x))$$

⋮

$$= \arg \min_{(g,h) \in \mathcal{G} \times \mathcal{H}} E_{\tilde{F}_x} \left[\frac{\|x - f(Z)\|^2}{2c} \right] + H(\tilde{F}_x, F_Z)$$

→ Now we have all three relevant distributions.



Efficient Encoding-Decoding Scheme

So far: For a given input x and a function f (Remember: $F_{X|Z} \sim \mathcal{N}(f(Z), cl)$) we can find an optimal approximation \tilde{F}_x^* .

Output generation process: For an input x sample $z \sim \tilde{F}_x^*$. Then sample the ouput $\hat{x} \sim F_{X|Z}(\cdot|z)$

Next: We want $\hat{x} = x$ with high probability
→ Optimize over $f \in \mathcal{F}$

$$f^* = \arg \max_{f \in \mathcal{F}} E_{\tilde{F}_x^*} [\log f_{X|Z}(x|Z)]$$

⋮

$$= \arg \min_{f \in \mathcal{F}} E_{\tilde{F}_x^*} \left[\frac{\|x - f(Z)\|^2}{2c} \right]$$



Optimization Problem

Combined optimization problem:

$$(f^*, g^*, h^*) = \arg \min_{(f, g, h) \in \mathcal{F} \times \mathcal{G} \times \mathcal{H}} E_{\tilde{F}_x^*} \left[\frac{\|x - f(Z)\|^2}{2c} \right] + H(\tilde{F}_x, F_Z)$$


Concrete implementation: Use neural networks for f, g, h and optimize with stochastic gradient descent.

Additional assumption: Covariance matrix of approximation \tilde{F}_x is diagonal
→ $h(x)$ is vector of diagonal entries




Reparametrisation Trick

So far:



→ No gradient for gradient descent

Now:






Loss Function

Minimize loss function:

$$\begin{aligned} & E_{\tilde{F}_x^*} \left[\frac{\|x - f(Z)\|^2}{2c} \right] + H(\tilde{F}_x, F_Z) \\ & \approx \frac{1}{2c} \|x - f(z)\|^2 + H(\mathcal{N}(g(x), h(x)), \mathcal{N}(0, I)) \end{aligned}$$

Here we used a Monte Carlo approximation over one sample for the expectation.

VAE Model



Application: Autoencoder [7]



Input				
PVAE				
VAE-123				
VAE-345				
<hr/>				
Input				
PVAE				
VAE-123				
VAE-345				



Application: Generator [7]

PVAE



VAE-123




VAE-345






Application: Linear Interpolation [7]

Latent vectors $z_{left}, z_{right} \rightarrow$ Interpolation $\alpha z_{left} + (1 - \alpha)z_{right}$, $\alpha \in [0, 1]$



MNIST example:





Application: Attribute Manipulation [7]

Calculate mean vectors $z_{pos_smiling}$, $z_{neg_smiling}$ and derive

$$z_{smiling} = z_{pos_smiling} - z_{neg_smiling} \rightarrow z_{left} + \alpha z_{smiling}, \alpha \in [0, 1]$$

$\alpha=0$ —————— $\alpha=1$



add
smiling
vector



subtract
smiling
vector



add
sunglass
vector



subtract
sunglass
vector



add
sunglass
vector

subtract
sunglass
vector

Generative Adversarial Network: A short story [3]




Generative Adversarial Network [6]



A GAN consists of two neural networks: a generator who produce new data and a discriminator who detect whether the data is fake or real.



Generative Adversarial Network





Loss Function [6, 10]

The generator and the discriminator play a minimax game. This can be expressed in the loss function

$$\min_G \max_D \text{Loss}(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Discriminator Loss:

$$\frac{1}{m} \sum_{i=1}^m [\log D(x^{(i)}) + \log(1 - D(G(z^{(i)})))]$$

Generator Loss:

$$\frac{1}{m} \sum_{i=1}^m [\log(1 - D(G(z^{(i)})))]$$

Instead of minimizing $\log(1 - D(G(z)))$ we can maximize $\log D(G(z))$
→ use Binary Cross-Entropy Loss



Train a GAN

```
1 from nn import MLP, Writer
2 from activations import Identity, Sigmoid, Tanh, ReLU, LeakyReLU, Softmax, Activation
3 from loss import MSE, CrossEntropy
4 from dataset import Dataset
5 from optimizer import Adam, SGD
6 from layers import Dense
7
8 # 1. Train discriminator
9 fake_img = self.generator.feedforward(self.sample(batchSize))
10 real_img = np.asarray(train[0])
11
12 input = np.concatenate((fake_img, real_img), axis = 1)
13 label = np.concatenate(
14     (
15         np.zeros((1, batchSize)),
16         0.9 * np.ones((1, batchSize))
17     ),
18     axis = 1
19 )
20
21 self.discriminator.feedforward(np.asarray(input))
22 self.discriminator.backpropagate(np.asarray(label), timeStep = i+1)
23 discriminatorLoss = self.discriminator.getLoss(np.asarray(label))
24
25 # 2. Train generator
26 self.discriminator.feedforward(fake_img)
27 self.discriminator.backpropagate(np.ones((1, batchSize)), timeStep = i+1,
28     updateParameters = False)
29 discriminatorGradient = self.discriminator.layers[0].gradient
30 self.generator.backpropagate(discriminatorGradient, timeStep = i+1,
31     useLoss = False)
32 generatorLoss = self.discriminator.getLoss(np.ones((1, batchSize)))
```

Problems of GANs



- highly unstable optimization problem
- mode collapse

What would the Seine at Argenteuil look like today?



What would the Seine at Argenteuil look like today?



Monet ↪ Photos



[13]

What would the Seine at Argenteuil look like today?



Monet ↪ Photos




photo → Monet

[13]




Application: CycleGAN [13]

- Image-to-image translation
- learn mapping $G : X \rightarrow Y$
- reverse mapping $F : Y \rightarrow X$ with $F(G(X)) \approx X$ and vice versa
→ Cycle Consistency





Application: CycleGAN [13]





Application: CycleGAN [13]

- mapping $G : X \rightarrow Y$

$$\begin{aligned}\mathcal{L}_{GAN}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]\end{aligned}$$

- similar for mapping $F : Y \rightarrow X$

$$\mathcal{L}_{GAN}(F, D_X, Y, X)$$

- cycle-consistency $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$ and
 $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$

$$\begin{aligned}\mathcal{L}_{cyc}(G, F) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1]\end{aligned}$$

→ full loss:

$$\begin{aligned}\mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{GAN}(G, D_Y, X, Y) \\ & + \mathcal{L}_{GAN}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{cyc}(G, F)\end{aligned}$$




Application: CycleGAN [13]





Application: CycleGAN [13]





Application: CycleGAN [13]






Table of Contents

1. Neural Networks

1.1 Activation Functions

1.2 Loss Functions

1.3 Optimizer

1.4 Code: Training a Classifier

2. Autoencoder

2.1 Principal Component Analysis

2.2 Autoencoder

2.3 Denoising Autoencoder

2.4 Variational Autoencoder

2.5 Generative Adversarial Network

3. Outlook

Outlook



- Convolutional Neural Networks
- Recurrent Neural Networks
- Residual Neural Networks
- Transfer Learning



Questions



Bibliography I

- [1] Andrew Ng (Stanford). *Lecture 14 — Machine Learning (Stanford)*. Youtube. 2008. URL:
<https://www.youtube.com/watch?v=ey2PE5xi9-A>.
- [2] Francois Chollet. *Building Autoencoders in Keras*. URL: <https://blog.keras.io/building-autoencoders-in-keras.html>.
- [3] deepart.io. URL: <https://deepart.io/#>.
- [4] Aurélien Géron. *Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems*. 1st. O'Reilly Media, Inc., 2017. ISBN: 978-1-491-96229-9.
- [5] Al Gharakhanian. LinkedIn. 2016. URL:
<https://www.linkedin.com/pulse/gans-one-hottest-topics-machine-learning-al-gharakhanian>.
- [6] Ian J. Goodfellow et al. *Generative Adversarial Networks*. 2014. arXiv: 1406.2661 [stat.ML].



Bibliography II

- [7] Xianxu Hou et al. *Deep Feature Consistent Variational Autoencoder*. 2016. arXiv: 1610.00291 [cs.CV].
- [8] Diederik P. Kingma and Jimmy Ba. *Adam: A Method for Stochastic Optimization*. 2014. arXiv: 1412.6980 [cs.LG].
- [9] Marie-Anne Lachaux et al. *Unsupervised Translation of Programming Languages*. 2020. arXiv: 2006.03511 [cs.CL].
- [10] Diego Gomez Mosquera. medium.com. URL:
<https://medium.com/ai-society/gans-from-scratch-1-a-deep-introduction-with-code-in-pytorch-and-tensorflow-cb03cdcdba0f>.
- [11] Michael A. Nielsen. *Neural Networks and Deep Learning*. Determination Press, 2015.
- [12] Jonathon Shlens. *A Tutorial on Principal Component Analysis*. 2014. arXiv: 1404.1100 [cs.LG].



Bibliography III

- [13] Jun-Yan Zhu et al. *Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks*. 2018. arXiv: 1703.10593 [cs.CV].